

Java-Success.com

Prepare to fast-track, choose & go places with 800+ Java & Big Data Q&As with lots of code & diagrams.

[Home](#) [Why? ▾](#) [300+ Java FAQs ▾](#) [300+ Big Data FAQs ▾](#) [Courses ▾](#)

[👤 Membership ▾](#) [Your Career ▾](#)

[Home](#) › [bigdata-success.com](#) › [Tutorials - Big Data](#) › [TUT - Starting Spark & Scala](#) ›

10. Setting up & getting started with Hive

10. Setting up & getting started with Hive

 Posted on [November 17, 2018](#)

Hive is popular SQL based tool to process Big Data. It stores & retrieves data to/from HDFS.

Prerequisite: Xcode & Hadoop are installed as outlined in [Setting up & getting started with Hadoop on Mac](#)

1. Install Hive

Let's install hive-2.3.4 with hadoop 2.7.7.

300+ Java Interview FAQs

300+ Java FAQs



16+ Java Key Areas Q&As



150+ Java Architect FAQs



80+ Java Code Quality Q&As



150+ Java Coding Q&As



300+ Big Data Interview FAQs

300+ Big Data FAQs



Tutorials - Big Data



TUT -  Starting Big Data

TUT - Starting Spark & Scala

```
1
2 $ wget http://www-us.apache.org/dist/hive/hive-2.3
3 $ sudo tar xvfz apache-hive-2.3.4-bin.tar.gz -C /usr
4
```

This will install 2.3.4 version of Hive in the folder
/usr/local/apache-hive-2.3.4-bin/

2. set HIVE_HOME

```
1
2 $ vi ~/.bash_profile
3
```

In the “`.bash_profile`”

```
1
2 export HIVE_HOME=/usr/local/apache-hive-2.3.4-bin
3 export HIVE_CONF_DIR=$HIVE_HOME/conf
4 export PATH=$HIVE_HOME/bin:$PATH
5 export CLASSPATH=$CLASSPATH:$HADOOP_HOME/lib/*
6 export CLASSPATH=$CLASSPATH:$HIVE_HOME/lib/*
7
```

Activate the change.

```
1
2 $ source ~/.bash_profile
3
```

3. Install the MySQL Server

Hive requires a RDBMS to store its meta-data. It is called the **Hive metastore**. Hive can be used with the embedded Derby database, but Derby is good for the sake of development and unit testing, but won't scale to a production environment as only a single user can connect to the derby database at any instant of time. Better way to configure is to use an external

TUT - Starting with Python

TUT - Kafka

TUT - Pig

TUT - Apache Storm

TUT - Spark Scala on Zeppelin

TUT - Cloudera

TUT - Cloudera on Docker

TUT - File Formats

TUT - Spark on Docker

TUT - Flume

TUT - Hadoop (HDFS)

TUT - HBase (NoSQL)

TUT - Hive (SQL)

TUT - Hadoop & Spark

TUT - MapReduce

TUT - Spark and Scala

TUT - Spark & Java

TUT - PySpark on Databricks

TUT - Zookeeper

800+ Java Interview Q&As

300+ Core Java Q&As



300+ Enterprise Java Q&As



150+ Java Frameworks Q&As



120+ Companion Tech Q&As



Tutorials - Enterprise Java



database which is JDBC compliant like MySQL, Oracle, etc.

```
1  
2 $ brew install mysql  
3
```

4. Start the MySQL Server

```
1  
2 $ mysql.server start  
3
```

5. Set up MySQL Server

```
1  
2 $ mysql -u root  
3
```

Create a database named “metastore” and a new user named “hiveuser”, and grant permissions. You also need to run the schema upgrade scripts. The Hive version being used here is “2.3.x”.

```
1  
2 mysql> CREATE DATABASE metastore;  
3 mysql> USE metastore;  
4 mysql> SOURCE /usr/local/apache-hive-2.3.4-bin/scripts/  
5 mysql> CREATE USER 'hiveuser'@'localhost' IDENTIFIED BY '  
6 mysql> GRANT SELECT,INSERT,UPDATE,DELETE,ALTER,CREATE  
7
```

Note: **schematool** is an offline command line tool to manage the **metastore**. This tool can be used to initialize the metastore schema for the current Hive version (E.g. 2.3.x).

```
1
2 $ schematool -dbType mssql -info
3
```

Before you run hive for the first time, run the following to initialize the schema:

```
1
2 $ schematool -dbType mssql -initSchema
3
```

6. Download mysql-connector-java

Download from

<https://dev.mysql.com/downloads/connector/j/>

```
1
2 $ curl -L 'http://www.mysql.com/get/Downloads/Connector-J/mysql-connector-java-8.0.13.jar'
3
```

Copy the “mysql-connector-java-8.0.13.jar” that has the Driver class “com.mysql.jdbc.Driver”

```
1
2 $ mv mysql-connector-java-8.0.13/mysql-connector-j
3
```

```
1
2 $ sudo chown -R <user>:admin apache-hive-2.3.4-bin
3
```

7. Configure Hive – hive-env.sh

```
1
2 $ cd /usr/local/apache-hive-2.3.4-bin/conf
3 $ cp hive-env.sh.template hive-env.sh
4
```

8. Configure Hive – hive-site.xml

```
1
2 $ cd /usr/local/apache-hive-2.3.4-bin/conf
3 $ cp hive-default.xml.template hive-site.xml
4
```

Make sure the following selected lines between the `<configuration>` and `</configuration>` tags of `hive-site.xml` have the values as shown below:

```
1
2 $ vi /usr/local/apache-hive-2.3.4-bin/conf/hive-site.xml
3
```

Make sure that the following properties are set as shown below. These properties are used to connect to the external MySQL metastore database. When you start the Hive shell, it will automatically connect to the MySQL database and create the required tables in it.

```
1
2 <configuration>
3
4 //.....
5
6 <property>
7   <name>javax.jdo.option.ConnectionPassword</name>
8   <value>password</value>
9   <description>password to use against metastore</description>
10 </property>
11
12 <property>
13   <name>javax.jdo.option.ConnectionURL</name>
14   <value>jdbc:mysql://localhost:3306/metastore?characterEncoding=utf8</value>
15   <description>JDBC connect string for a JDBC metastore</description>
16 </property>
17
18 <property>
19   <name>javax.jdo.option.ConnectionDriverName</name>
20   <value>com.mysql.cj.jdbc.Driver</value>
21   <description>Driver class name for a JDBC metastore</description>
22 </property>
23
24 <name>javax.jdo.option.ConnectionUserName</name>
```

```

25     <value>hiveuser</value>
26     <description>Username to use against metastore</description>
27 </property>
28
29 <property>
30     <name>hive.querylog.location</name>
31     <value>/usr/local/apache-hive-2.3.4-bin/iotmp</value>
32     <description>Location of Hive run time structure</description>
33 </property>
34
35 <property>
36     <name>hive.exec.local.scratchdir</name>
37     <value>/usr/local/apache-hive-2.3.4-bin/iotmp</value>
38     <description>Local scratch space for Hive jobs</description>
39 </property>
40 <property>
41     <name>hive.downloaded.resources.dir</name>
42     <value>/usr/local/apache-hive-2.3.4-bin/iotmp</value>
43     <description>Temporary local directory for added</description>
44     </description>
45 </property>
46 <property>
47     <name>hive.metastore.schema.verification</name>
48     <value>false</value>
49     <description>
50         Enforce metastore schema version consistency
51         True: Verify that version information stored in the
52             schema migration attempt. Users are responsible for
53             proper metastore schema migration. (Default)
54         False: Warn if the version information stored in the
55     </description>
56 </property>
57
58 //.....
59
60 <configuration>
61

```

Include the below configuration in **conf/hive-site.xml**

```

1 <property>
2     <name>hive.metastore.uris</name>
3     <value>thrift://localhost:9083</value>
4 </property>
5

```

9. Configure YARN – yarn-site.xml

Make sure that the following properties are set in yarn-site.xml.

```
1  ...
2  <property>
3      <name>yarn.nodemanager.aux-services</name>
4      <value>mapreduce_shuffle</value>
5  </property>
6
7  <property>
8      <name>yarn.nodemanager.aux-services.mapreduce_
9      <value>org.apache.hadoop.mapred.ShuffleHandler
10 </property>
11 ....
12
```

10. Start hadoop if not already started

```
1
2 $ ./start-all.sh
3 $ jps -lm
4
```

Output:

```
1 4580 sun.tools.jps.Jps -lm
2 4406 org.apache.hadoop.yarn.server.resourcemanager.ResourceManager
3 4167 org.apache.hadoop.hdfs.server.datanode.DataNode
4 4503 org.apache.hadoop.yarn.server.nodemanager.NodeManager
5 4073 org.apache.hadoop.hdfs.server.namenode.NameNode
6 4283 org.apache.hadoop.hdfs.server.namenode.SecondaryNameNode
7
```

11. Verify Hive installation

```
1
2 $ echo $HIVE_HOME
3
```

Output:

```
1  
2 /usr/local/apache-hive-2.3.4-bin  
3
```

12. Start the Hive metastore interface

```
1 $ hive --service metastore &
```

The Hive metastore interface by default listens at port 9083. Make sure it is.

```
1 $ netstat -an | grep 9083
```

13. Start Hive

```
1  
2 $ hive  
3
```

Output:

```
1  
2 hive>  
3
```

14. Create HDFS path

```
1  
2 $ hdfs dfs -mkdir -p /user/hive/warehouse  
3 $ hdfs dfs -chown -R hiveuser /user/hive/warehouse  
4 $ hdfs dfs -chmod -R 777 /user/hive/warehouse  
5
```

15. Create a Hive table

```
1  
2 hive> CREATE TABLE IF NOT EXISTS test_table
```



```
3 > (col1 int COMMENT 'Integer Column',
4 > col2 string COMMENT 'String Column')
5 > COMMENT 'This is test table'
6 > ROW FORMAT DELIMITED
7 > FIELDS TERMINATED BY ','
8 > STORED AS TEXTFILE;
9 OK
10 Time taken: 4.016 seconds
11
```

16. Insert data into test_table

```
1
2 hive> INSERT INTO default.test_table values(1, 'Some test value')
3 hive> INSERT INTO test_table values(2, 'Some test value')
4
```

Output: runs a mapreduce job

```
1 Total jobs = 3
2 Launching Job 1 out of 3
3 Number of reduce tasks is set to 0 since there's no reducer in this job.
4 Starting Job = job_1542456661191_0002, Tracking URL = http://localhost:8080/jobdetails.jsp?jobid=job_1542456661191_0002
5 Kill Command = /usr/local/hadoop-2.7.7/bin/hadoop-kill
6 Hadoop job information for Stage-1: number of maps=0, number of reducers=0
7 2018-11-17 23:15:50,403 Stage-1 map = 0%, reduce = 0%
8 2018-11-17 23:15:55,562 Stage-1 map = 100%, reduce = 0%
9 Ended Job = job_1542456661191_0002
10 Stage-4 is selected by condition resolver.
11 Stage-3 is filtered out by condition resolver.
12 Stage-5 is filtered out by condition resolver.
13 Moving data to directory hdfs://localhost:9000/user/hive/warehouse/default.db/test_table
14 Loading data to table default.test_table
15 MapReduce Jobs Launched:
16 Stage-Stage-1: Map: 1 HDFS Read: 4344 HDFS Write: 1024
17 Total MapReduce CPU Time Spent: 0 msec
18 OK
19 Time taken: 13.153 seconds
20
```

17. SELECT from test_table

```
1
2 hive> SELECT * FROM test_table;
3 OK
4 1 Some test value 1
```

```

5 | 2      Some test value 2
6 | Time taken: 0.118 seconds, Fetched: 2 row(s)
7 |

```

18. Where is the underlying data stored in HDFS?

```

1 | hdfs dfs -ls /user/hive/warehouse

```

Output:

```

1 | drwxrwxrwx  - user supergroup      0 2018-11-1

```

```

1 | hdfs dfs -ls /user/hive/warehouse/test_table

```

Output:

```

1 | -rwxrwxrwx  3 user supergroup      20 2018-11-1
2 | -rwxrwxrwx  3 user supergroup      20 2018-11-1

```

```

1 | hdfs dfs -cat /user/hive/warehouse/test_table/00000

```

Output:

```

1 | 1,Some test value 1

```

```

1 | hdfs dfs -cat /user/hive/warehouse/test_table/00000

```

Output:

```

1 | 2,Some test value 2

```

So, basically when you execute a SQL query, a mapreduce job is run to save or select data to/from

HDFS. Hive can be run with other execution engines like Spark & Tez. The default engine is mapreduce.

We will look at more examples in the coming tutorials.

Learn more about Hive Q&As style

Hive Interview FAQs

◀ 09. Getting started with Spark & Hadoop – client mode on local & cluster mode on yarn

13 Apache Sqoop interview Q&As ▶

Disclaimer

The contents in this Java-Success are copyrighted and from EmpoweringTech pty ltd. The EmpoweringTech pty ltd has the right to correct or enhance the current content without any prior notice. These are general advice only, and one needs to take his/her own circumstances into consideration. The EmpoweringTech pty ltd will not be held liable for any damages caused or alleged to be caused either directly or indirectly by these materials and resources. Any trademarked names or labels used in this blog remain the property of their respective trademark owners. Links to external sites do not imply endorsement of the linked-to sites. [Privacy Policy](#).