

Java-Success.com

Prepare to fast-track, choose & go places with 800+ Java & Big Data Q&As with lots of code & diagrams.

[Home](#) [Why? ▾](#) [300+ Java FAQs ▾](#) [300+ Big Data FAQs ▾](#) [Courses ▾](#)


[👤 Membership ▾](#) [Your Career ▾](#)

[Home](#) > [bigdata-success.com](#) > [Tutorials - Big Data](#) > [TUT - Cloudera on Docker](#) > 10:

Docker Tutorial: Hadoop Big Data services & folders on Cloudera quickstart

10: Docker Tutorial: Hadoop Big Data services & folders on Cloudera quickstart

 Posted on [May 26, 2019](#)

You can also install it on VMWare as illustrated on the  [Getting started with BigData on Cloudera](#).

If you are not familiar with Docker get some hands-on experience at [a series of step by step Docker tutorials](#) with Java & Springboot examples.

This tutorial is based on [09: Docker Tutorial: Cloudera on Docker via DockerHub](#), where Cloudera

300+ Java Interview FAQs

300+ Java FAQs



16+ Java Key Areas Q&As



150+ Java Architect FAQs



80+ Java Code Quality Q&As



150+ Java Coding Q&As



300+ Big Data Interview FAQs

300+ Big Data FAQs



Tutorials - Big Data



TUT -  Starting Big Data

TUT - Starting Spark & Scala

Quickstart gets installed on Docker for learning purpose.

```
1 ~/projects/docker-hadoop]$ docker run --hostname=qu
2 --privileged=true -t -i -v /Users/arulkumarankuma
3 p 8888:8888 -p 80:80 -p 7180:7180 cloudera/quicksto
```

When the above command runs, you will see all the services that gets started:

zookeeper, journalnode, datanode, namenode, and secondary namenode

```
1 Starting zookeeper ... STARTED
2 starting datanode, logging to /var/log/hadoop-hdfs,
3 Started Hadoop datanode (hadoop-hdfs-datanode):
4 starting journalnode, logging to /var/log/hadoop-h
5 Started Hadoop journalnode: [ OK ]
6 starting namenode, logging to /var/log/hadoop-hdfs,
7 Started Hadoop namenode:
8 starting secondarynamenode, logging to /var/log/ha
9 Started Hadoop secondarynamenode:
```

historyserver, nodemanager, resourcemanager, HBase master, rest, thrift, Hive Metastore, Hive Server2, and Sqoop Server

```
1 Started Hadoop historyserver:
2 starting nodemanager, logging to /var/log/hadoop-y
3 Started Hadoop nodemanager:
4 starting resourcemanager, logging to /var/log/had
5 Started Hadoop resourcemanager:
6 starting master, logging to /var/log/hbase/hbase-l
7 Started HBase master daemon (hbase-master):
8 starting rest, logging to /var/log/hbase/hbase-hb
9 Started HBase rest daemon (hbase-rest):
10 starting thrift, logging to /var/log/hbase/hbase-l
11 Started HBase thrift daemon (hbase-thrift):
12 Starting Hive Metastore (hive-metastore):
13 Started Hive Server2 (hive-server2):
14 Starting Sqoop Server:
15 Sqoop home directory: /usr/lib/sqoop2
16 Setting SQOOP_HTTP_PORT: 12000
```

TUT - Starting with Python

TUT - Kafka

TUT - Pig

TUT - Apache Storm

TUT - Spark Scala on Zeppelin

TUT - Cloudera

TUT - Cloudera on Docker

TUT - File Formats

TUT - Spark on Docker

TUT - Flume

TUT - Hadoop (HDFS)

TUT - HBase (NoSQL)

TUT - Hive (SQL)

TUT - Hadoop & Spark

TUT - MapReduce

TUT - Spark and Scala

TUT - Spark & Java

TUT - PySpark on Databricks

TUT - Zookeeper

800+ Java Interview Q&As

300+ Core Java Q&As



300+ Enterprise Java Q&As



150+ Java Frameworks Q&As



120+ Companion Tech Q&As



Tutorials - Enterprise Java



```
17 | Setting SQOOP_ADMIN_PORT: 12001
```

Spark history-server, HBase regionserver, hue, and Impala

```
1 | Starting Spark history-server (spark-history-server):
2 | Starting Hadoop HBase regionserver daemon: starting
3 | hbase-regionserver.
4 | Starting hue:
5 | Started Impala State Store Server (statestored):
6 |
7 | Setting OOZIE_HOME: /usr/lib/oozie
8 | Sourcing: /usr/lib/oozie/bin/oozie.sh
9 | setting JAVA_LIBRARY_PATH="$JAVA_LIBRARY_PATH:/usr/lib/oozie/lib"
10 |
11 | .....
12 | Starting Solr server daemon:
13 | Started Impala Catalog Server (catalogd) :
14 | Started Impala Server (impalad):
15 |
```

All these services can be viewed via the “Cloudera Manager” admin console.

/etc/init.d services

init.d is the sub-directory of **/etc** directory in Linux file system. init.d basically contains the bunch of start/stop/reload/restart/status scripts which are used to control the Hadoop ecosystem daemons whilst the system is running or during boot. If you look at /etc/init.d then you will notice all the scripts for different services

```
1 | [root@quickstart /]# cd /etc/init.d
2 | [root@quickstart init.d]# ls
3 | atd                                hadoop-httpfs
4 | cloudera-quickstart-init           hadoop-mapreduce-h
5 | cloudera-scm-agent                 hadoop-yarn-nodemar
6 | cloudera-scm-server                hadoop-yarn-proxysc
7 | crond                              hadoop-yarn-resourc
8 | flume-ng-agent                     halt
```

```

9  functions                                hbase-master
10  hadoop-hdfs-datanode                    hbase-regionserver
11  hadoop-hdfs-journalnode                 hbase-rest
12  hadoop-hdfs-namenode                    hbase-solr-indexer
13  hadoop-hdfs-secondarynamenode          hbase-thrift
14  [root@quickstart init.d]# service hbase-master start
15  HBase master daemon is running
16  [root@quickstart init.d]# service mysqld status
17  mysqld (pid 169) is running...
18  [root@quickstart init.d]# service impala-server start
19  Impala Server is running
20

```

netstat -anp to find ports

mysql runs on port 3306.

```

1  [root@quickstart init.d]# netstat -anp | grep mysq
2  tcp        0      0 0.0.0.0:3306          0.0.0.0:*        LISTENING
3  tcp        0      0 127.0.0.1:3306       127.0.0.1:*      LISTENING
4  tcp        0      0 172.17.0.2:3306     172.17.0.2:*    LISTENING
5  tcp        0      0 172.17.0.2:3306     172.17.0.2:*    LISTENING
6  tcp        0      0 127.0.0.1:3306       127.0.0.1:*      LISTENING
7  tcp        0      0 127.0.0.1:3306       127.0.0.1:*      LISTENING
8  tcp        0      0 127.0.0.1:3306       127.0.0.1:*      LISTENING
9  tcp        0      0 172.17.0.2:3306     172.17.0.2:*    LISTENING
10 tcp        0      0 172.17.0.2:3306     172.17.0.2:*    LISTENING
11 unix  2      [ ACC ]     STREAM  LISTENING
12 [root@quickstart init.d]#
13

```

impalad runs on multiple ports, and **21050** for **impalad front-end**, 21000 for impalad impala-shell, 22000 is for back-end, and so on. You can check the Cloudera documentation for further details.

```

1  [root@quickstart init.d]# netstat -anp | grep impa
2  tcp        0      0 0.0.0.0:23000        0.0.0.0:*        LISTENING
3  tcp        0      0 0.0.0.0:21050       0.0.0.0:*        LISTENING
4  tcp        0      0 0.0.0.0:21000       0.0.0.0:*        LISTENING
5  tcp        0      0 0.0.0.0:25000       0.0.0.0:*        LISTENING
6  tcp        0      0 0.0.0.0:22000       0.0.0.0:*        LISTENING
7  tcp        0      0 172.17.0.2:23000    172.17.0.2:*    LISTENING
8  tcp        0      0 172.17.0.2:23000    172.17.0.2:*    LISTENING
9  tcp        0      0 127.0.0.1:45732     127.0.0.1:*      LISTENING
10 unix  3      [ ]       STREAM  CONNECTED

```

```

11 | unix  3      [ ]      STREAM  CONNECTED
12 | unix  3      [ ]      STREAM  CONNECTED
13 | unix  3      [ ]      STREAM  CONNECTED
14 | unix  3      [ ]      STREAM  CONNECTED
15 | unix  3      [ ]      STREAM  CONNECTED
16 | unix  3      [ ]      STREAM  CONNECTED
17 | unix  3      [ ]      STREAM  CONNECTED
18 | unix  3      [ ]      STREAM  CONNECTED
19 | unix  2      [ ]      STREAM  CONNECTED
20 |

```

ps auxwww to find service run details

```

1 [root@quickstart init.d]# ps auxwww | grep hive-ser
2 hive      2084  0.4  2.4 793748 234952 ?        S1
3 r/log/hive -Dhive.log.file=hive-server2.log -Dhive
4 .log -Dhadoop.home.dir=/usr/lib/hadoop -Dhadoop.id
5 tive -Dhadoop.policy.file=hadoop-policy.xml -Djava
6 doop.util.RunJar /usr/lib/hive/lib/hive-service-1.1
7 root      5616  0.0  0.0 103300 1988 pts/0    S+
8 [root@quickstart init.d]#
9

```

Hbase starts a number of services like master, region server, etc.

```

1 [root@quickstart init.d]# ps auxwww | grep hbase
2 hbase      1329  0.0  0.0  9256  2424 ?        S
3 f foreground_start master
4 hbase      1343  0.5  2.0 3145200 198992 ?        S1
5 moryError=kill -9 %p -XX:+UseConcMarkSweepGC -XX:R
6 hbase-hbase-master-quickstart.cloudera.log -Dhbase
7 library.path=/usr/lib/hadoop/lib/native:/usr/lib/
8 .hbase.master.HMaster start
9 hbase      1472  0.0  0.0  9256  2396 ?        S
10 f foreground_start rest
11 hbase      1486  0.2  1.3 2955260 128120 ?        S1
12 ryError=kill -9 %p -XX:+UseConcMarkSweepGC -Dhbase
13 -Dhbase.home.dir=/usr/lib/hbase -Dhbase.id.str=hba
14 r/lib/hbase/lib/native/Linux-amd64-64 -Dhbase.secu
15 hbase      1695  0.0  0.0  9256  2452 ?        S
16 f foreground_start thrift
17 hbase      1711  0.2  1.5 2963200 148116 ?        S1
18 moryError=kill -9 %p -XX:+UseConcMarkSweepGC -Dhba
19 log -Dhbase.home.dir=/usr/lib/hbase -Dhbase.id.sti
20 :/usr/lib/hbase/lib/native/Linux-amd64-64 -Dhbase
21 hbase      2659  0.0  0.0  9256  2420 ?        S
22 f foreground_start regionserver

```

```

23 hbase      2673  0.5  1.9 3129228 189668 ?      SL
24 utOfMemoryError=kill -9 %p -XX:+UseConcMarkSweepGC
25 root      5644  0.0  0.0 103300  2032 pts/0    S+
26 [root@quickstart init.d]#
27

```

/var/log folder

This where the log files go.

```

1 [root@quickstart /]# cd /var/log
2 [root@quickstart log]# ls
3 btmap          flume-ng          hadoop
4 cloudera-scm-agent  hadoop-0.20-mapreduce  hadoop
5 cloudera-scm-server  hadoop-hdfs          hbase
6 dedup.log        hadoop-httpfs        hbase
7 dracut.log        hadoop-kms           hive
8 [root@quickstart log]# cd hive
9 [root@quickstart hive]# ls -ltr
10 total 16
11 -rw-r--r--  1 hive hive  198 May 26 01:38 hive-metacat
12 -rw-r--r--  1 hive hive  167 May 26 01:38 hive-service
13 -rw-r--r--  1 hive hive  439 May 26 01:38 hive-metacat
14 -rw-r--r--  1 hive hive 1128 May 26 01:39 hive-service
15 [root@quickstart hive]#
16

```

Java & Python versions

```

1 [root@quickstart /]# java -version
2 java version "1.7.0_67"
3 Java(TM) SE Runtime Environment (build 1.7.0_67-b01)
4 Java HotSpot(TM) 64-Bit Server VM (build 24.65-b04)
5 [root@quickstart /]# which java
6 /usr/bin/java
7 [root@quickstart /]# python --version
8 Python 2.6.6
9 [root@quickstart cloudera]# which python
10 /usr/bin/python
11

```

/usr/bin folder

```

1 [root@quickstart lib]# ls -ltr /usr/bin | grep imp
2 -rwxr-xr-x  1 root root  1856 Mar 23 2016 imp

```

```

3 -rwxr-xr-x  1 root root    11030 Mar 23  2016 im
4 [root@quickstart lib]#
5

```

/usr/lib folder

```

1 [root@quickstart lib]# ls -ltr /usr/lib
2 ....
3 [root@quickstart lib]# ls -ltr /usr/lib/hive/lib
4 .....
5 lrwxrwxrwx  1 root root    38 Apr  6  2016 accum
6 lrwxrwxrwx  1 root root    37 Apr  6  2016 accum
7 lrwxrwxrwx  1 root root    37 Apr  6  2016 accum
8 lrwxrwxrwx  1 root root    40 Apr  6  2016 mysql
9

```

/usr/jars folder

All the jars used above in “/usr/lib/...”

```

1 [root@quickstart lib]# ls -ltr /usr/jars

```

/var/run/ or /run folder

Run-time variable data. You can get the “pid” (i.e process id). You will also know what services are running.

```

1 [root@quickstart cloudera]# ls -ltr /var/run/hadoop
2 total 16
3 -rw-r--r--  1 hdfs hdfs 4 May 26 09:08 hadoop-hdfs-c
4 srw-rw-rw-  1 hdfs hdfs 0 May 26 09:08 dn.50010
5 -rw-r--r--  1 hdfs hdfs 4 May 26 09:08 hadoop-hdfs-
6 -rw-r--r--  1 hdfs hdfs 4 May 26 09:09 hadoop-hdfs-
7 -rw-r--r--  1 hdfs hdfs 4 May 26 09:09 hadoop-hdfs-
8

```

```

1 [root@quickstart cloudera]# ls -ltr /var/run/impala
2 total 12
3 -rw-r--r--  1 impala impala 5 May 26 09:10 statestor
4 -rw-r--r--  1 impala impala 5 May 26 09:10 catalogd
5 -rw-r--r--  1 impala impala 5 May 26 09:10 impalad-

```

```
6 [root@quickstart cloudera]#  
7
```

Examples from Cloudera quickstart

The jar shown below has a number of examples, and you can test your environment by running the MapReduce job as shown.

```
1 [root@quickstart /]# find / -name hadoop-mapreduce-  
2 /usr/lib/hadoop-mapreduce/hadoop-mapreduce-example
```

The jobs that are available in hadoop-mapreduce-example.jar:

```
1 .....  
2     pgd.addClass("wordcount", WordCount.class,  
3                 "A map/reduce program that coun  
4     pgd.addClass("wordmean", WordMean.class,  
5                 "A map/reduce program that coun  
6     pgd.addClass("wordmedian", WordMedian.class,  
7                 "A map/reduce program that coun  
8     pgd.addClass("wordstandarddeviation", WordSt  
9                 "A map/reduce program that coun  
10    pgd.addClass("aggregatwordcount", Aggregate  
11                "An Aggregate based map/reduce  
12    pgd.addClass("aggregatwordhist", Aggregate  
13                "An Aggregate based map/reduce  
14    pgd.addClass("grep", Grep.class,  
15                "A map/reduce program that coun  
16    pgd.addClass("randomwriter", RandomWriter.c  
17                "A map/reduce program that writ  
18    pgd.addClass("randomtextwriter", RandomText  
19                "A map/reduce program that writes 10GB of re  
20    pgd.addClass("sort", Sort.class, "A map/redu  
21  
22    pgd.addClass("pi", QuasiMonteCarlo.class, Qu  
23    pgd.addClass("bbp", BaileyBorweinPlouffe.cl  
24    pgd.addClass("distbbp", DistBbp.class, Dist  
25  
26    pgd.addClass("pentomino", DistributedPentom  
27                "A map/reduce tile laying program to find so  
28    pgd.addClass("secondarysort", SecondarySort  
29                "An example defining a secondar  
30    pgd.addClass("sudoku", Sudoku.class, "A sude  
31    pgd.addClass("join", Join.class, "A job tha
```



```

32     pgd.addClass("multifilewc", MultiFileWordCo
33     pgd.addClass("dbcount", DBCountPageView.clas
34     pgd.addClass("teragen", TeraGen.class, "Gene
35     pgd.addClass("terasort", TeraSort.class, "Ru
36     pgd.addClass("teravalidate", TeraValidate.c
37     exitCode = pgd.run(argv);
38     .....
39

```

Run a mapreduce job

Running the “pi” example **MapReduce** job:

```

1 [root@quickstart /]# sudo -u hdfs hadoop jar \
2 > /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examp
3 > pi 10 100
4 Number of Maps = 10
5 Samples per Map = 100
6 Wrote input for Map #0
7 Wrote input for Map #1
8 Wrote input for Map #2
9 Wrote input for Map #3
10 Wrote input for Map #4
11 Wrote input for Map #5
12 Wrote input for Map #6
13 Wrote input for Map #7
14 Wrote input for Map #8
15 Wrote input for Map #9
16 Starting Job
17 19/05/26 03:27:00 INFO client.RMPProxy: Connecting
18 19/05/26 03:27:01 INFO input.FileInputFormat: Total
19 19/05/26 03:27:01 INFO mapreduce.JobSubmitter: num
20 19/05/26 03:27:01 INFO mapreduce.JobSubmitter: Sub
21 19/05/26 03:27:01 INFO impl.YarnClientImpl: Submit
22 19/05/26 03:27:01 INFO mapreduce.Job: The url to t
23 /
24 19/05/26 03:27:01 INFO mapreduce.Job: Running job
25 19/05/26 03:27:08 INFO mapreduce.Job: Job job_1558
26 19/05/26 03:27:08 INFO mapreduce.Job: map 0% redu
27 19/05/26 03:27:14 INFO mapreduce.Job: map 10% red
28 19/05/26 03:27:15 INFO mapreduce.Job: map 20% red
29 19/05/26 03:27:16 INFO mapreduce.Job: map 30% red
30 19/05/26 03:27:18 INFO mapreduce.Job: map 40% red
31 19/05/26 03:27:19 INFO mapreduce.Job: map 50% red
32 19/05/26 03:27:20 INFO mapreduce.Job: map 60% red
33 19/05/26 03:27:21 INFO mapreduce.Job: map 70% red
34 19/05/26 03:27:23 INFO mapreduce.Job: map 80% red
35 19/05/26 03:27:24 INFO mapreduce.Job: map 100% re
36 19/05/26 03:27:25 INFO mapreduce.Job: map 100% re
37 19/05/26 03:27:25 INFO mapreduce.Job: Job job_1558
38 19/05/26 03:27:25 INFO mapreduce.Job: Counters: 49

```

File System Counters

```
FILE: Number of bytes read=226
FILE: Number of bytes written=1259
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=2730
HDFS: Number of bytes written=215
HDFS: Number of read operations=43
HDFS: Number of large read operations=0
HDFS: Number of write operations=3
```

Job Counters

```
Launched map tasks=10
Launched reduce tasks=1
Data-local map tasks=10
Total time spent by all maps in org.apache.hadoop.mapred.MapTask=100
Total time spent by all reduces in org.apache.hadoop.mapred.ReduceTask=100
Total time spent by all map tasks=100
Total time spent by all reduce tasks=100
Total vcore-seconds taken by all map tasks=100
Total vcore-seconds taken by all reduce tasks=100
Total megabyte-seconds taken by all map tasks=100
Total megabyte-seconds taken by all reduce tasks=100
```

Map-Reduce Framework

```
Map input records=10
Map output records=20
Map output bytes=180
Map output materialized bytes=280
Input split bytes=1550
Combine input records=0
Combine output records=0
Reduce input groups=2
Reduce shuffle bytes=280
Reduce input records=20
Reduce output records=0
Spilled Records=40
Shuffled Maps =10
Failed Shuffles=0
Merged Map outputs=10
GC time elapsed (ms)=308
CPU time spent (ms)=4690
Physical memory (bytes) snapshot=3000000
Virtual memory (bytes) snapshot=1000000
Total committed heap usage (bytes)=1000000
```

Shuffle Errors

```
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
```

File Input Format Counters

```
Bytes Read=1180
```

File Output Format Counters

```
Bytes Written=97
```

```
94 Job Finished in 24.374 seconds
95 Estimated value of Pi is 3.14800000000000000000
96 [root@quickstart /]#
97
```

Do you want to open multiple terminal windows?

```
1 $ docker ps
2 CONTAINER ID          IMAGE                COMMAND
3 bf645c6a2930         cloudera/quickstart  "/usr/bin/
4
1 $ docker exec -it bf645c6a2930 /bin/bash
2 [root@quickstart /]#
3
```

How to stop the container?

```
1 $ docker ps
2 CONTAINER ID          IMAGE                COMMAND
3 bf645c6a2930         cloudera/quickstart  "/usr/bin/
4
1 $ docker stop bf645c6a2930
2
```

What is next?

In the next post let's look at the [CLIs like impala-shell, hdfs, hive, spark-shell, pyspark](#), etc.

◀ 09: Docker Tutorial: Getting started with Hadoop Big Data on Cloudera quickstart

11: Docker Tutorial: Hadoop Big Data CLIs on Cloudera quickstart ▶

Disclaimer

The contents in this Java-Success are copyrighted and from EmpoweringTech pty ltd. The EmpoweringTech pty ltd has the right to correct or enhance the current content without any prior notice. These are general advice only, and one needs to take his/her own circumstances into consideration. The EmpoweringTech pty ltd will not be held liable for any damages caused or alleged to be caused either directly or indirectly by these materials and resources. Any trademarked names or labels used in this blog remain the property of their respective trademark owners. Links to external sites do not imply endorsement of the linked-to sites. [Privacy Policy](#)

© 2022 [java-success.com](https://www.java-success.com)