

Java-Success.com

Prepare to fast-track, choose & go places with 800+ Java & Big Data Q&As with lots of code & diagrams.

[Home](#) [Why? ▾](#) [300+ Java FAQs ▾](#) [300+ Big Data FAQs ▾](#) [Courses ▾](#)

[Membership ▾](#) [Your Career ▾](#)

[Home](#) > [bigdata-success.com](#) > [Tutorials - Big Data](#) > [TUT - Cloudera on Docker](#) > 20:

Docker Tutorial: Apache Spark (spark-submit) in Java on Cloudera quickstart

20: Docker Tutorial: Apache Spark (spark-submit) in Java on Cloudera quickstart

 Posted on [June 7, 2019](#)

This extends [Docker Tutorial: BigData on Cloudera quickstart via Docker](#).

Step 1: Run the container on a command line.

```
1 ~/projects/docker-hadoop]$ docker run --hostname=q
2 --privileged=true -t -i -v /Users/arulkumarankumar
3 --publish-all=true -p 8888:8888 -p 80:80 -p 7180:71
```

300+ Java Interview FAQs

300+ Java FAQs



16+ Java Key Areas Q&As



150+ Java Architect FAQs



80+ Java Code Quality Q&As



150+ Java Coding Q&As



300+ Big Data Interview FAQs

300+ Big Data FAQs



[Tutorials - Big Data](#)



[TUT - Starting Big Data](#)

[TUT - Starting Spark & Scala](#)

Step 2: Java is already installed as part of the “cloudera/quickstart” docker image.

```
1 [root@quickstart /]# java -version
2 java version "1.7.0_67"
3 Java(TM) SE Runtime Environment (build 1.7.0_67-b01)
4 Java HotSpot(TM) 64-Bit Server VM (build 24.65-b04)
5
```

update-alternatives

If you have multiple versions or installations of Java, you can list them as shown below:

```
1 [root@quickstart /]# update-alternatives --config java
```

Install maven

Step 3: Download maven from “<http://mirror.ventraip.net.au/apache/maven/maven-3/3.6.1/binaries/apache-maven-3.6.1-bin.tar.gz>”

```
1 [root@quickstart /]# sudo curl -L -O http://mirror.ventraip.net.au/apache/maven/maven-3/3.6.1/binaries/apache-maven-3.6.1-bin.tar.gz
2
```

```
1 [root@quickstart /]# cp apache-maven-3.6.1-bin.tar.gz /opt
2 [root@quickstart /]# cd /opt
3 [root@quickstart /]# sudo tar xvfz apache-maven-3.6.1-bin.tar.gz
4 [root@quickstart /]# sudo rm -f apache-maven-3.6.1-bin.tar.gz
5
```

Step 4: Add it to your profile so that you can run the “mvn” command. Add it to “.bashrc” if you are a root user or to “.bash_profile” if you are a any other user.

```
1 [root@quickstart opt]# vi ~/.bash_profile
```

TUT - Starting with Python

TUT - Kafka

TUT - Pig

TUT - Apache Storm

TUT - Spark Scala on Zeppelin

TUT - Cloudera

TUT - Cloudera on Docker

TUT - File Formats

TUT - Spark on Docker

TUT - Flume

TUT - Hadoop (HDFS)

TUT - HBase (NoSQL)

TUT - Hive (SQL)

TUT - Hadoop & Spark

TUT - MapReduce

TUT - Spark and Scala

TUT - Spark & Java

TUT - PySpark on Databricks

TUT - Zookeeper

800+ Java Interview Q&As

300+ Core Java Q&As



300+ Enterprise Java Q&As



150+ Java Frameworks Q&As



120+ Companion Tech Q&As



Tutorials - Enterprise Java



add the following lines at the end of “~/.bash_profile”

```
1 M3_HOME=/opt/apache-maven-3.6.1
2 export MAVEN_OPTS=-Dhttps.protocols=TLSv1,TLSv1.1,
3 export PATH=$PATH:$M3_HOME/bin
4
```

Apply the changes with the “source” command.

```
1 [root@quickstart opt]# source ~/.bash_profile
```

```
1 [root@quickstart opt]# mvn -version
2 Apache Maven 3.6.1 (d66c9c0b3152b2e69ee9bac180bb8f
3 Maven home: /opt/apache-maven-3.6.1
4 Java version: 1.7.0_67, vendor: Oracle Corporation
5 Default locale: en_US, platform encoding: UTF-8
6 OS name: "linux", version: "4.9.125-linuxkit", arch
7 [root@quickstart opt]#
8
```

Now Maven is installed.

Note: Even though the maven is already in the docker image we have installed the latest version to a different location “/opt/”.

```
1 [root@quickstart /]# ls -ltr /usr/local/apache-mav
2 total 4
3 drwxr-xr-x 6 root root 4096 Apr  6 2016 apache-mav
4 [root@quickstart /]#
```

Create a simple maven project structure

Step 5: Create “projects” folder and create a maven project structure.

```
1 [root@quickstart ~]# cd /
```

```
2 [root@quickstart ~]# mkdir /projects
3 [root@quickstart ~]# cd /projects
4
```

Create a maven project structure.

```
1 [root@quickstart projects]# mvn archetype:generate
2   -DartifactId=my-app \
3   -DarchetypeArtifactId=maven-archetype-quickstart
4   -DinteractiveMode=false
5
```

Give it a few minutes.

```
1 [root@quickstart projects]# tree my-app
2 my-app
3 |— pom.xml
4 |— src
5 |   |— main
6 |   |   |— java
7 |   |   |   |— com
8 |   |   |   |   |— mycompany
9 |   |   |   |   |   |— app
10 |   |   |   |   |       |— App.java
11 |   |— test
12 |   |   |— java
13 |   |   |   |— com
14 |   |   |   |   |— mycompany
15 |   |   |   |   |   |— app
16 |   |   |   |   |       |— AppTest.java
17
18 11 directories, 3 files
19 [root@quickstart projects]#
20
```

Create a simple Spark job in Java

Step 6: Add the Spark library to maven pom.xml file.

```
1 [root@quickstart projects]# cd my-app
2 [root@quickstart my-app]# vi pom.xml
3
```

The pom.xml file should look like

```
1 <project xmlns="http://maven.apache.org/POM/4.0.0"
2   xsi:schemaLocation="http://maven.apache.org/POM
3   <modelVersion>4.0.0</modelVersion>
4   <groupId>com.mycompany.app</groupId>
5   <artifactId>my-app</artifactId>
6   <packaging>jar</packaging>
7   <version>1.0-SNAPSHOT</version>
8   <name>my-app</name>
9   <url>http://maven.apache.org</url>
10  <dependencies>
11
12    <dependency>
13      <groupId>junit</groupId>
14      <artifactId>junit</artifactId>
15      <version>3.8.1</version>
16      <scope>test</scope>
17    </dependency>
18    <!-- https://mvnrepository.com/artifact/org.ap
19    <dependency>
20      <groupId>org.apache.spark</groupId>
21      <artifactId>spark-core_2.10</artifactId>
22      <version>1.6.0</version>
23    </dependency>
24
25  </dependencies>
26 </project>
27
```

Step 7: Create a simple Spark job in Java.

```
1 [root@quickstart my-app]# vi src/main/java/com/myco
```

The “SimpleSpark.java” file looks like:

```
1 package com.mycompany.app;
2
3 import org.apache.spark.api.java.*;
4 import org.apache.spark.SparkConf;
5 import java.util.*;
6 import org.apache.spark.api.java.function.Function
7
8 public class SimpleSpark {
9
10     public static void main (String[] args) {
11         SparkConf conf = new SparkConf().setAppNam
12         JavaSparkContext sc = new JavaSparkContext
13
```

```
14     List<String> data = Arrays.asList("John",
15     JavaRDD<String> rdd = sc.parallelize(data);
16     System.out.println(rdd.collect());
17
18 }
19
20 }
21
```

compile & package it as a jar file with maven

Step 8: Compile & package it as a jar file.

```
1 [root@quickstart my-app]# mvn clean package
```

```
1 [root@quickstart my-app]# tree
2 .
3 |__ pom.xml
4 |__ src
5 |   |__ main
6 |       |__ java
7 |           |__ com
8 |               |__ mycompany
9 |                   |__ app
10 |                       |__ App.java
11 |                       |__ SimpleSpark.java
12 |   |__ test
13 |       |__ java
14 |           |__ com
15 |               |__ mycompany
16 |                   |__ app
17 |                       |__ AppTest.java
18 |__ target
19 |   |__ classes
20 |       |__ com
21 |           |__ mycompany
22 |               |__ app
23 |                   |__ App.class
24 |                   |__ SimpleSpark.class
25 |   |__ maven-archiver
26 |       |__ pom.properties
27 |   |__ maven-status
28 |       |__ maven-compiler-plugin
29 |           |__ compile
30 |               |__ default-compile
31 |                   |__ createdFiles.lst
32 |                   |__ inputFiles.lst
33 |           |__ testCompile
34 |               |__ default-testCompile
```

```
35 |         |         | createdFiles.lst
36 |         |         | inputFiles.lst
37 |         | my-app-1.0-SNAPSHOT.jar
38 |         | surefire-reports
39 |         | | com.mycompany.app.AppTest.txt
40 |         | | TEST-com.mycompany.app.AppTest.xml
41 |         | test-classes
42 |         | | com
43 |         | | | mycompany
44 |         | | | | app
45 |         | | | | AppTest.class
46 |
47 | 28 directories, 15 files
48 | [root@quickstart my-app]#
```

“my-app-1.0-SNAPSHOT.jar” is the packaged jar file in the “target” folder.

spark-submit to run the spark job

Step 9: Run the Spark job in the jar file via Spark-submit command.

Local client mode

```
1 [root@quickstart my-app]# spark-submit \
2 --class com.mycompany.app.SimpleSpark \
3 --master local \
4 --deploy-mode client \
5 target/my-app-1.0-SNAPSHOT.jar
6
```

```
1 ...
2 [John, Peter, Samuel]
3 ....
4
```

Cluster yarn mode

```
1 [root@quickstart my-app]# spark-submit \
2 --class com.mycompany.app.SimpleSpark \
3 --master yarn \
4 --deploy-mode cluster \
5 target/my-app-1.0-SNAPSHOT.jar
6
```

The job will be submitted to yarn and executed. In order to look at the logs you need to check the logs for the printed values.

◀ 19: Docker Tutorial: Apache Spark SQL – on Cloudera quickstart

21: Docker Tutorial: Apache Spark (spark-submit) in Scala on Cloudera quickstart ▶

Disclaimer

The contents in this Java-Success are copyrighted and from EmpoweringTech pty Ltd. The EmpoweringTech pty Ltd has the right to correct or enhance the current content without any prior notice. These are general advice only, and one needs to take his/her own circumstances into consideration. The EmpoweringTech pty Ltd will not be held liable for any damages caused or alleged to be caused either directly or indirectly by these materials and resources. Any trademarked names or labels used in this blog remain the property of their respective trademark owners. Links to external sites do not imply endorsement of the linked-to sites. [Privacy Policy](#)